

EinOps Whoopsydoopsy

```
# Matrix multiplication (2D x 2D)
c = einsum('ij,jk->ik', a, b)

# Batch matrix multiplication (3D x 3D)
c = einsum('bij,bjk->bik', a, b)


# Dot product
c = einsum('i,i->', a, b)

# Outer product
c = einsum('i,j->ij', a, b)
```

$$M_{ij} = \sum_k A_{ik} B_{kj} = A_{ik} B_{kj}$$


Implementing your own Neural Network

Explanation of building blocks



nn.Module

- The base class for all neural network components
- Allows you to organize layers and their parameters modularly
- Automatically tracks all parameters for optimization
- Must implement forward() method defining computation

nn.Parameter

- Special tensor that gets automatically tracked
- Used for learnable weights and biases
- Will be updated during backpropagation

- Example: `self.weight = nn.Parameter(torch.randn(3, 4))`

Buffers:

- tracked by the model, part of statedict
- not updated by the optimizer

- Example: `self.register_buffer('running_mean', torch.zeros(hidden_dim))`

```
self.running_mean = 0.9 * self.running_mean + 0.1 * batch_mean
```

nn.Sequential



- Container to stack modules in sequence
- Automatically chains forward() calls
- Useful for linear architectures

- Example:

```
nn.Sequential(  
    nn.Linear(10, 20),  
    nn.ReLU(),  
    nn.Linear(20, 30)  
)
```


Multilayer Perceptron

- The classic simplest neural network
- A stack of fully connected layers with activation functions
- No sense of space in the neural network structure
- Works with any kind of data
- Also, still part of Transformer architecture



Convolution

- You can make use of the “local” structure of the image
- An “edge detector” (for example) acts the same on parts of the image




Parameters

```
kernal: 3  
stride: 1  
padding: 0  
output_padding: 0  
dilation: 1
```

■ Kernel
■ Output
■ Input

Convolution - stride

- You can give the image, as it passes through the network, lower resolution, by skipping steps
- Here, no pixel gets completely lost




Parameters

```
kernal: 3  
stride: 2  
padding: 0  
output_padding: 0  
dilation: 1
```

■ Kernel
■ Output
■ Input

Convolution - padding

- Pixels on the border of the image get viewed only once, image gets smaller
- Image gets smaller without stride
- Add “empty” pixels outside the image to cancel this effect



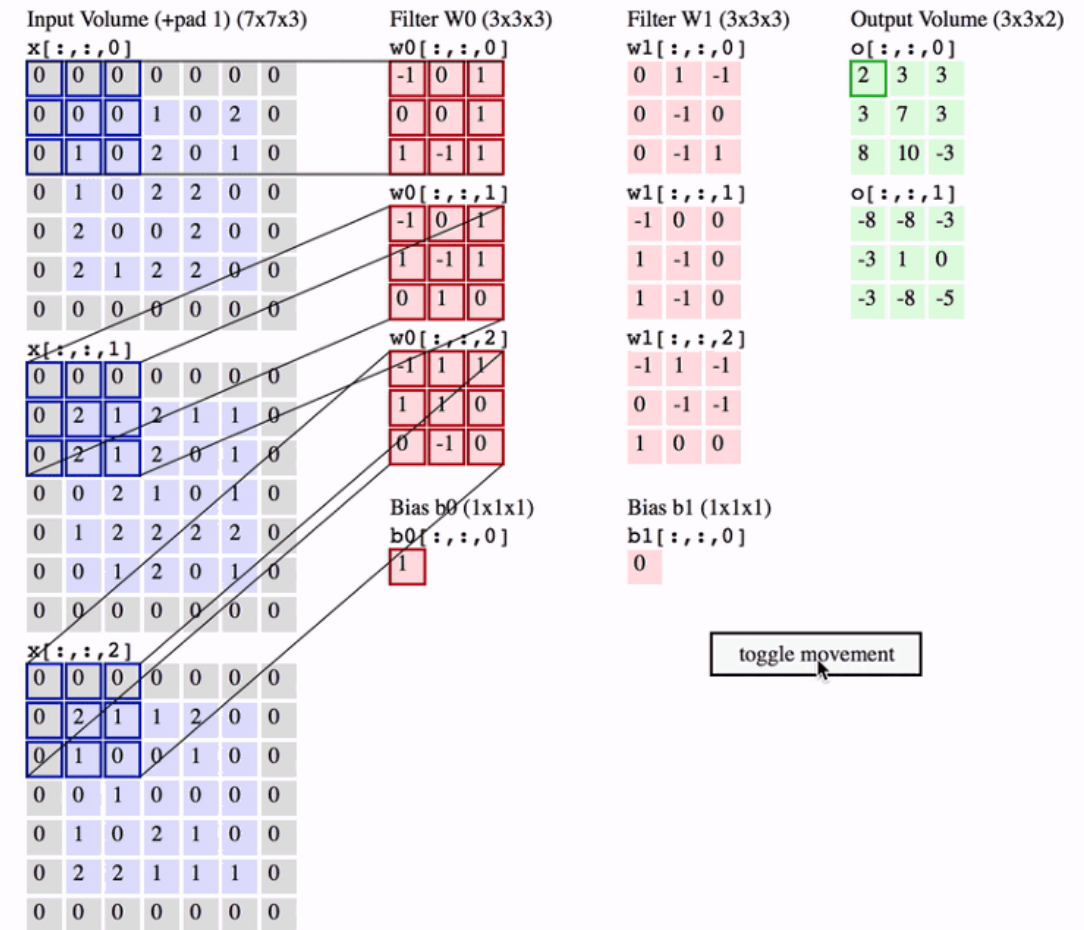
Parameters

```
kernal: 3  
stride: 2  
padding: 1  
output_padding: 0  
dilation: 1
```

- Padding
- Kernel
- Output
- Input



Convolution - channels

- Typically, you have more than one convolution kernel active at once
- This leads to an additional output dimension per picture: channel
- Example: Colour or different features like different edges/objects



Batch Norm

- The value of the activation is less important than its relative size
- Example: brightness in an image / across images
- Normalize distribution



Batch Norm– how this gets calculated

Learnable parameters

$$y = \frac{x - E[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$$


In Training: Updated as a rolling average
In Deployment: Frozen

The diagram shows the Batch Normalization formula with several annotations. Two blue arrows point from the text 'Learnable parameters' to the parameters γ and β in the formula. Two orange arrows point from the text 'In Training: Updated as a rolling average' and 'In Deployment: Frozen' to the variance term $\text{Var}[x]$. Two green arrows point from the text 'Activations: For taking activation statistics, all activations from the same channel are taken as draws from the same distribution (independent of location in picture)' to the input x and the mean term $E[x]$.

Activations: For taking activation statistics, all activations from the same channel are taken as draws from the same distribution (independent of location in picture)

MaxPooling (2d)

- Reduce the size of the resulting Image
- Taking the maximum of each kernel window
- Example: Object detection in a part of the image



AveragePool

- Reduce resulting image to a single number
- Example: classifying an image with a single label
- Average over space dimensions

ResNet 34

